

Chapter 9: Filesystems

Introduction

Till this point, we have seen how a kernel of an operating system works as a resource manager. Given that both the processor and the main memory are resources in the system, 539kernel manages these resources ¹ and provides them to the different processes in the system. Another role of a kernel is to provide a way to communicate with external devices, such as the keyboard and hard disk. *Device drivers* are the way of realizing this role of the kernel ². The details of the external devices and how to communicate with them are low-level and may be changed at any time. The goal of a device driver is to communicate with a given device by using the device's own language and the other goal of a device driver is to provide an interface for any other component of the system that wish to use the given device, most probably the low-level details of this given device will be hidden behind the interface that the device driver provides, that means the user of the device drivers doesn't need to know anything about how the device really work.

The matter of hiding the low-level details with something higher-level is too important and can be found in, basically, everywhere in computing and the kernels are not an exception of that. Of course, there is virtually no limit of providing higher-level concepts based a previous lower-level concept, also upon something that we consider as a high-level concept we can build something even higher-level. Beside the previous example of device drivers, one of obvious example where the kernels fulfill the role of hiding the low-level details and providing something higher-level, in other words, providing an *abstraction*, is a filesystem which provides the well-known abstraction, a file.

In this chapter we are going to cover these two topics, device drivers and filesystem by using 539kernel. As you may recall, it turned out that accessing to the hard disk is an important aspect for virtual memory, so, to be able to implement virtual memory, the kernel itself needs to access the hard disk which makes it an important component in the kernel, so, we are going to implement a device driver that communicate with the hard disk in this chapter. After getting the ability of reading from the hard disk or writing to it, we can explore the idea of providing abstractions by the kernel through writing a filesystem that uses the hard disk device driver and provides a higher-level view of the hard disk, that we all familiar with, instead of the physical view of the hard disk which has been described previously in chapter . The final result of this chapter is version NE of 539kernel which has as we mentioned a hard disk device driver and a filesystem.

¹Incompletely of course, to keep 539kernel as simple as possible, only the basic parts of resources management were presented.

²At least a monolithic kernel.

ATA Device Driver

No need to say the hard disks are too common devices that are used as secondary storage devices. Also, there are a lot of manufacturers that manufacture hard disks and sell them. Imagine for a moment that each hard disk from a different manufacturer use its own way for the communication, that is, the method `X` should be used to be able to communicate with hard disks from manufacturer `A` while the method `Y` should be used with hard disks from manufacturer `B`. Given that there are too many manufacturers, this will be a nightmare. Each hard disk will need its own device driver which talks a different language from the other hard disk device drivers. Fortunately, this is not the case, at least for the hard disks, in this type of situations, standards are here to the rescue. A manufacturer may design the hard disk hardware in anyway, but when it comes to the part of the communication between the hard disk and the outside world, a standard can be used, so, any device driver that works with this given standard, will be able to communicate with this new hard disk. There are many well-known standards that are related to the hard disks, *small computer system interface* (SCSI) is one of them, another one is *advanced technology attachment* (ATA)³. While the latter one is more common in personal computers we are going to focus on it here and write a device driver for it, the former one is more common in servers.

As in PIC which is been discussed in chapter , ATA hard disks can be communicated with by using port-mapped I/O communication through the instructions `in` and `out` that we have covered previously. But before discussing the ATA commands that let us to issue a read or write request to the hard disk, let's write two routines in assembly that can be used in C code and perform the same functionality for the instructions `in` and `out`. If you didn't recall, the instruction `out` is used to write some bytes to a given port number, so, if we know the port number that a device receives that commands from, we can use the instruction `out` to write a valid command to that port, on the other hand, the instruction `in` reads data from a given port, for example, sometimes after we send a command to a device, it responds through writing something on a specific port, the instruction `in` can be used to read this value. The assembly code of the both routines that we are going to define next should reside in `starter.asm` anywhere between `bits 32` and the beginning of `start_kernel` routine. As we have said previously, the goal of these two routines is to make it possible for C code to use the instructions `in` and `out` through calling these routines. The following is the code of `dev_write` which can be used by C kernel code to write to a given port. In C, we can see that it has this prototype: `dev_write(int port, int cmd)`.

```
dev_write:
    ; Part 1
```

³Another well-known name for ATA is *Integrated Drive Electronics* (IDE). The older ATA standard is now known as Parallel ATA (PATA) while the newer version of ATA is known is Serial ATA (SATA).

```

push edx
push eax

; Part 2
xor edx, edx
xor eax, eax

; Part 3
mov dx, [esp + 12]
mov al, [esp + 16]

; Part 4
out dx, al

; Part 5
pop eax
pop edx

ret

```

The core part of this routine is part four which contains the instruction `out` that send the value which is stored in `al` to the port number which is stored in `dx`. Because we are using these two registers ⁴, we push their previous values into the stack as we did in the first part of the routine, pushing the previous values of these registers lets us restore them easily after the routine finishes its work, this restoration is performed in the fifth part of the routine right before returning from the routine, this is an important step to make sure that when the routine returns, the environment of the caller is same as the one before calling the routine. After storing the previous values of `eax` and `edx` we can use them freely, so, the first step after that is to clear their previous values by setting the value 0 to the both of them, as you can see, we have used `xor` and the both operands of it are the same register that we wish to clear, this is a well-known way in assembly programming to clear the value of a register ⁵. After that, we can move the values that have been passed to the routine as parameters to the correct registers to be used with `out` instruction, this is performed in the third part of the routine ⁶. Also, we need to define another routine called `dev_write_word` which is exactly same as `dev_write` but write a word (2 bytes) instead of one byte to a port. The following is the code of this routine.

```

dev_write_word:
    push edx
    push eax

```

⁴Which are as you know parts of the registers `eax` and `edx` respectively.

⁵To my best knowledge its more performant than the normal way of using `mov`.

⁶The readers who have previous knowledge in x86 assembly programming may notice that I've omitted the epilogue of routines which creates a new stack frame, this decision has been made to make the matters simpler and you are absolutely free to use the calling convention.

```

xor edx, edx
xor eax, eax

mov dx, [esp + 12]
mov ax, [esp + 16]

out dx, ax

pop eax
pop edx

ret

```

As you can see, the only difference between `dev_write` and `dev_write_word` is that the first one uses the register `al` (8 bit) as the second operand of `out` while the second one uses `ax` (16 bit) instead, so, a word can be written to the port. The following is the code of the routine `dev_read` which uses the instruction `in` to read the data from a given port and return them to the caller, its prototype can be imagined as `char dev_read(int port)`.

```

dev_read:
    push edx

    xor edx, edx
    xor eax, eax

    mov dx, [esp + 8]

    in ax, dx

    pop edx

    ret

```

For the same reason of restoring the previous environment when returning to the caller, the routine pushes the value of `edx` into the stack, then both of `edx` and `eax` are cleared since they will be used by the instruction `in`. After that, the value of the passed parameter which represents the port number that caller wishes to read from, is stored in `dx`. Finally, `in` is called, the result is stored in `ax`⁷, the previous value of `edx` is restored and the routine returns. You may ask, why did we only stored and restored the previous value of `edx` while the register `eax` is also used also, why didn't we store and restore the previous value of `eax`?

⁷Since the first operand of `in` is `ax` and not `al` then a **word** will be read from the port and not a single byte. The decision on using `ax` instead of `al` was made here because of our needs as you will see later, if you need to read just one byte for some reason you can define another routine for that.

The reason is that `dev_read` is a function that returns a value, and according to `cdecl` convention the returned values should be stored in the register in `eax`, so, the value of `eax` is intended to be changed when return to the caller, therefore, it will not be correct, logically, to restore the the previous value of `eax` when `dev_read` returns. The ultimate goal of defining both `dev_write` and `dev_read` is to make them available to be used in C code, so, the lines `global dev_write`, `global dev_write_word` and `global dev_read` should be written in the beginning of `starter.asm`.

The Driver

One ATA bus in the computer's motherboard makes it possible to attach two hard disks into the system, one of them is called master drive which is the main one that the computer boots from, the other disk is known as slave drive. Usually, a computer comes with two ATA buses instead of just one, which means up two four hard disks can be attached into the computer. The first one of those buses is known as the primary bus while the second one is known as the secondary bus. The port numbers that are used to communicate with the devices that are attached into the primary bus start from `0x1F0` and ends in `0x1F7` each one of them has its own functionality while the ports number from `0x170` to `0x177` are used to communicate with devices that are attached into the secondary bus, so, there are eight ports for each ATA bus. The terms that combine a bus name and a device name are used to specify exactly which device is being discussed, for example, primary master means the master hard disk that is connected to the primary bus while secondary slave means the slave hard disk which is connected to the secondary bus. For the sake of simplicity, our device driver is going to assume that there is only a primary master and all read and write requests should be oriented to this primary master, therefore, our device driver uses the port number `0x1F0` as the base port to send the commands via PIC. You may ask, why are we calling this port number a base port? As you know that all the following port numbers are valid to communicate with the primary ATA bus: `0x1F0`, `0x1F1`, `0x1F2`, `0x1F3`, `0x1F4`, `0x1F5`, `0x1F6`, `0x1F7`, we can add any numbers from 0 through 7 to the base port number of the primary bus `0x1F0` to get a correct port number to communicate with the primary bus, the same is correct with the secondary ATA bus which its base port number is `0x170`. So, we can define the base port as a macro ⁸ as we will see in our device driver, and the we can use this macro by adding a specific value to it from 0 through 7 to get a specific port, the advantage of doing so is the easiness of changing the value of the base port to another port without the need of changing the code itself. Before starting in the implementation of the driver, let's create new two files: `ata.h` and `ata.c` which will contain the code of the device driver which communicates to ATA devices and provides an interface for the rest of the kernel to write and read data from the disk. The following is the content of `ata.h` and

⁸Or even variable.

the details of the functions will be discussed in the next subsections.

```
#define BASE_PORT 0x1F0
#define SECTOR_SIZE 512

void *read_disk( int );
void write_disk( int, short * );

void *read_disk_chs( int );
void write_disk_chs( int, short * );
```

Addressing Mode

As in the main memory, the hard disks use addresses to read the data that are stored in a specific area of the disk, the same is applicable in write operation, the same address can be used to write on the same specific area. There are two schemes of hard disk addresses, the older one is known as *cylinder-head-sector* addressing (CHS) while the newer one which more dominant now is known as *logical block addressing* (LBA). In chapter we have covered the physical structure of hard disks and we know from that discussion that the data are stored in small blocks known as sectors, also, there are tracks which each one of them consists of a number of sectors, and finally, there are heads that should be positioned on a specific sector to read from it or to write to it. The scheme CHS uses the same concepts of physical structure of hard disk, the address of given sector on the hard disk can be composed by combining three numbers together, the cylinder (track) that this sector reside on, the sector that we would like to access and the head that is able to access this sector, however, this scheme is obsolete now and LBA is used instead of it. In LBA, a logical view of a hard disk is used instead of the physical view. This logical view states that the hard disk is composed of a number of logical blocks with a fixed size, say, n bytes. These blocks are contiguous in a similar way of the main memory, to reach any block you can use its own address and the addresses start from 0, the block right after the first one has the address 1 and so on. As you can see, addressing in LBA is more like the addressing of the main memory, the main difference here is that in current computers each address of the main memory points to a byte in memory while an address in LBA points to a block which can be a sector (512 bytes) or even bigger.

Reading from Disk

In this subsection we are going to implements both `read_disk_chs` and `read_disk` which send commands to an ATA hard disk via the available ports in order to read a sector/block from the disk. The first one of those functions uses CHS scheme while the second one uses LBA scheme. In the next discussions, I'm going to use the symbol `base_port` to indicate the base port of one of ATA ports, in our case, as we mentioned earlier, the base port is 0x1F0 since we are

going to use the primary bus in our device driver, but what we are discussing is applicable to any ATA bus with any base port number.

To issue a read command the command 0x20 should be sent to `base_port + 7`, but before doing that, a number of values should be set in the other ports in order to specify that address that we would like to read from, therefore, these values should be set before issuing the read command. In `base_port + 2` the number of sectors/blocks that we would like to read in this operation should be set. The value that should be written to `base_port + 6` specifies more than one thing, bit 6 of this value specifies whether we are using CHS or LBA in the current read request, when the bit's value is 0 then CHS is being used while the value 1 means that LBA is being used. The bits 5 and 7 of this value should always be 1. The bit 4 is used to specify the drive that we would like to read from, the value 0 for the master drive while 1 for the slave drive. In the case that we are using CHS, then the first four bits (0 to 3) of this value is used to specify the head and in LBA the bits 24 to 27 of the address should be stored in these bits. When the current addressing mode is CHS, the sector number that we would like our read operation to start from should be sent to `base_port + 3`, the low part of the cylinder number should be sent to `base_port + 4` and the high part of the cylinder number should be sent to `base_port + 5`. Once the read command is issued with the right parameters passed to the correct ports, we can read the value of `base_port + 7` to check if the disk finished the reading operation or not by reading the eighth bit (bit 7) of the value, when the value of this bit is 1 that means the drive is busy, once it becomes 0 that means that the operation completed. When the reading operation is completed successfully, the data are brought to `base_port` which means we need to read from it and put the required data in the main memory. The following is the code of `read_disk_chs` that should reside in `ata.c`⁹.

```
void *read_disk_chs( int sector )
{
    // Part 1

    dev_write( BASE_PORT + 6, 0x0a0 );
    dev_write( BASE_PORT + 2, 1 );
    dev_write( BASE_PORT + 3, sector );
    dev_write( BASE_PORT + 4, 0 );
    dev_write( BASE_PORT + 5, 0 );
    dev_write( BASE_PORT + 7, 0x20 );

    // ... //

    // Part 2

    int status = 0;
```

⁹Don't forget to include `ata.h` when you create `ata.c`

```

do
{
    status = dev_read( BASE_PORT + 7 );
} while ( ( status ^ 0x80 ) == 128 );

// ... //

// Part 3

short *buffer = kalloc( SECTOR_SIZE );

for ( int currByte = 0; currByte < ( SECTOR_SIZE / 2 ); currByte++ )
    buffer[ currByte ] = dev_read( BASE_PORT );

return buffer;
}

```

In the first part of `read_disk_chs` we send the required values to the appropriate ports as we have described above. In the port `base_port + 6` we set that the drive is 0 and the head is 0, also, we set that the addressing mode that we are using is CHS. In port `base + 2` we set that we would like to read one sector, that is, the function `read_disk_chs` reads 512 bytes from disk with each call. In port `base_port + 3` we set the number of the sector that we would like to read, as you can see, this number is passed through the parameter `sector`, so, the caller can specify the sector that it would like to read. In both `base_port + 4` and `base_port + 5` we specify the cylinder that we would like to read from, which is cylinder 0. Finally, we issue a read request to the ATA bus by writing 0x20 to `base_port + 7`. As you can see, for the sake of simplicity, we have used fixed values for a number of parameters here, in real situations, those fixed values should be more flexible. However, using LBA will provide us more flexibility with the simplicity that we are interested on. After issuing a read request in `read_disk_chs`, the second part of the function starts to work, it reads the value of the port `base_port + 7` which is going to contain the status of the drive, as we mentioned earlier, the eighth bit (bit 7) of this value indicates whether the drive is busy or not, with each iteration of the loop in part 2 the function reads this value and checks whether the value of the eighth bit is 1 by using bitwise operations, if this is the case that means that the drive is still busy in completing the reading operation, so, we need to wait for it until it finishes and we spend this time of waiting in reading the value of the port and check the eighth bits over and over again until the drive finishes, this technique of checking the device's status over and over again until it finishes an operation is known as *busy waiting*. When the read operation finishes, we can read the target content word by word ^{10 11} from the base port and that's the job of the third

¹⁰Recall that a "word" in x86 architecture is 2 bytes.

¹¹The fact that we need to read a word here instead of a byte is the reason of using the

part of `read_disk_chs` which reads a word each time and stores it in a buffer that we dynamically allocate from the kernel's heap. We have used the type `short` for the variable `buffer` because the size of this type in 32-bit architecture is 2 bytes, that is, a word. In the condition of the `for` loop we used `currByte < (SECTOR_SIZE / 2)` instead of `currByte < SECTOR_SIZE` due to the fact that we are reading two bytes in each iteration instead of one byte. Finally, the memory address which `buffer` points to, is going to contain the sector that we have just read from the disk, this memory address will be returned to the caller.

Now, we can turn to the other read function `read_disk` which uses LBA scheme instead of CHS. In fact, this new function will be almost similar to `read_disk_chs`, the main differences will be with some values that we are passing to the ports of ATA bus. In LBA scheme, `base_port + 6` should be changed to indicate that the scheme that we are using is LBA and we can do that as we mentioned before by setting the value 1 to bit 6. The other difference in this port value is that the bits 0 to 3 should contain the bits 24 to 27 of the logical block address that we would like to read from, the other parts of the addresses are divided to the following ports: `base_port + 3` contains bits 0 to 7 of that address, `base_port + 4` contains bits 8 to 15, `base_port + 5` contains bits 16 to 23. Both ports `base_port + 2` and `base_port + 7` stay the same. The function `read_disk` receives a parameter names `address` instead of `sector`, that is, the logical block address that the caller would like to read the data from should be passed to the function, by using bitwise operations the value of this parameter can be divided into the described parts to be filled in the appropriate ports. The rest of `read_disk` is exactly same as `read_disk_chs`. To not get a lot of space, the following is the beginning of the new function and the only part that has differences with `read_disk_chs`.

```
void *read_disk( int address )
{
    dev_write( BASE_PORT + 6, ( 0x0e0 | ( ( address & 0x0F000000 ) >> 24 ) ) );
    dev_write( BASE_PORT + 2, 1 );
    dev_write( BASE_PORT + 3, address & 0x000000FF );
    dev_write( BASE_PORT + 4, ( address & 0x0000FF00 ) >> 8 );
    dev_write( BASE_PORT + 5, ( address & 0x00FF0000 ) >> 16 );
    dev_write( BASE_PORT + 7, 0x20 );
}
```

Writing to Disk

In both CHS and LBA, write operation is called via ports exactly in the same way of reading, though, there are two differences. First, the command number to issue a write request is 0x30 which should be written to `base_port + 7` after setting the correct values to the port. Second, after the drive becomes ready to write the required data, we should write a word after the other to `base_port` until we finish, this is performed by calling the routine `dev_write_word` which

register `ax` instead of `al` as the first operand for `in` instruction in the function `dev_read`.

uses the instruction out to perform its job. Need no more discussion, the following two functions are `write_disk_chs` which uses CHS to write to the disk, and `write_disk` which uses LBA to write to the disk. Both of them receive a parameter called `buffer` which is a pointer to the data that we would like to write to the disk.

```
void write_disk_chs( int sector, short *buffer )
{
    dev_write( BASE_PORT + 6, 0x0a0 );
    dev_write( BASE_PORT + 2, 1 );
    dev_write( BASE_PORT + 3, sector );
    dev_write( BASE_PORT + 4, 0 );
    dev_write( BASE_PORT + 5, 0 );
    dev_write( BASE_PORT + 7, 0x30 );

    int status = 0;

    do
    {
        status = dev_read( BASE_PORT + 7 );
    } while ( ( status ^ 0x80 ) == 128 );

    for ( int currByte = 0; currByte < ( SECTOR_SIZE / 2 ); currByte++ )
        dev_write_word( BASE_PORT, buffer[ currByte ] );
}

void write_disk( int address, short *buffer )
{
    dev_write( BASE_PORT + 6, ( 0x0e0 | ( ( address & 0x0F000000 ) >> 24 ) ) );
    dev_write( BASE_PORT + 2, 1 );
    dev_write( BASE_PORT + 3, address & 0x000000FF );
    dev_write( BASE_PORT + 4, ( address & 0x0000FF00 ) >> 8 );
    dev_write( BASE_PORT + 5, ( address & 0x00FF0000 ) >> 16 );
    dev_write( BASE_PORT + 7, 0x30 );

    int status = 0;

    do
    {
        status = dev_read( BASE_PORT + 7 );
    } while ( ( status ^ 0x80 ) == 128 );

    for ( int currByte = 0; currByte < ( SECTOR_SIZE / 2 ); currByte++ )
        dev_write_word( BASE_PORT, buffer[ currByte ] );
}
```

Filesystem

Until now, we have seen multiple examples of presenting a logical higher-level view of something that is low-level, another example of that is a filesystem. A filesystem is a higher-level view of a storage device, this higher-level view makes it easier for the human user to use storage devices. Take hard disks as an example any imagine that we use them based on their low-level view that we have discussed in this chapter, it will be really to hard for us human to remember what is the logical block address of some data that we need to fetch right now, also, it will be really harder to remember which blocks of the hard disk are free and which of them are not in order to store new data. As alternative, we can build a logical view based on this low-level method of hard disk's functionality to make it easier for the human user to use the hard disk and hide the low-level complicated details of hard disks work.

Filesystems provide a well-known abstraction called *file* which is a sequence of bytes that has a name and is reachable by using this name. By providing this abstraction, it will be really easy for the human or even the applications to use the hard disk to store and retrieve data. Beside file abstraction, a filesystem may provide more features, for example, directories (or folders) are usually provided by filesystems in order to let the use to organize her files in a hierarchical way. To be able to implement the file abstraction, a way of maintaining the list of current files in the system and their locations in the disk should be available, due to that, specific filesystems usually use some sort of data structures to maintain the information about the files in a given system, that is, the *metadata* which means information about information. This data structure is stored in the disk for the future use and is interpretable by the kernel that implements that filesystem, by loading this data structure from the disk and interpreting it correctly, you can reach to the files and their organization as the user of the system created them.

The meaning of the term filesystem may differ based on the context. The first meaning of this term which we meant above is a component of an operating system that works based on a specific design to manage the files and directories of a system user. The other meaning that you may encounter more often in system administration context is the list of all files and directories that the users created in a specific system. In our next discussions I'm going to use the term filesystem to mean the first definition, while I'll use *run-time filesystem* to mean the second definition. Also, when the terms address is used in the next discussions, it means a logical block address.

As in programming languages, a filesystem may be divided into two parts: A design and an implementation. The design of a filesystem tells us how this filesystem stores the information about the run-time filesystem, how these metadata data are organized, which data structure is used to fulfill the goal of the filesystem an so on ¹². Of course, the design can be there, written on the papers as a brilliant theoretical piece of work. To realize this design, an

¹²The design part of programming languages is the language specifications.

implementation should be written that uses this design ¹³. For example, a well-known filesystem is *FAT* (short for: file allocation table) which is an old filesystem that started on 1977 and still in use nowadays. Because its design is available, anyone can write an implementation of it and make her kernel able to read run-time filesystems that have been created by another operating system that uses FAT, Linux kernel is an example of the kernels that have a FAT implementation. As an example of filesystem, we are going to design and implement *539filesystem* in this section.

The Design of 539filesystem

To make the implementation of 539filesystem as simple as possible, many restrictions have been presented in its design. First, there is only one directory in the run-time filesystem that uses 539filesystem which is the root directory, all the files that are created by the user are stored in this directory and there is no way to create and nested directories. Second, the maximum size of a file is 512 bytes and finally, when a file is created there is no way to modify it, its content is written when its created and that's it! I know, you may think that these are too harsh restrictions, while I agree, but these restrictions help in providing a too simple and elegant implementation of a real filesystem that can be extended easily to remove these restrictions.

The data structure that will be used to maintain the run-time filesystem and store it to the disk is linked-list which is a simple data structure that is known for its slow search operation but fast insert operation, its simplicity is the reason of choosing it for 539filesystem, real modern filesystems use more complex data structures to make the performance of the filesystem better. The block that has the address 100 in the disk is known as the *base block* in 539filesystem, from this block we can reach the whole run-time filesystem. The base block is divided into two parts, the size of each one of them is 4 bytes, the first one if the address of the metadata of the first file that has been created in the run-time filesystem, that is, the *head file* ¹⁴ while the second part if the address of the metadata of the last file that has been created in the run-time filesystem, that is, the *tail file* ¹⁵. Each file has its own metadata that contains file's name and the *next field* which stores the metadata address of the next file that has been created. The length of the filename is 256 bytes and the size of the next field is 4 bytes. It should be obvious now how can we reach all files in a run-time filesystem that uses 539filesystem, starting from the base block we get the metadata address of the head file and by using the next field we can reach the metadata of the next file and the process continues until we reach the tail file. When there is no next file, the value 0 is stored in the next field of the last file's metadata, that is, the tail file. The metadata of each file is stored in the block right before the content

¹³The implementation part of programming languages is a compiler or interpreter.

¹⁴In linked-list data structure's term: the head

¹⁵In linked-list data structure's term: the tail

of file which will be stored in one block only given that the size of a block is 512 bytes ¹⁶. For example, if the metadata data of file A is stored in the address 103, then the content of this file is stored in the address 104. By using this design, the basic functionalities of filesystems can be provided.

The Implementation of 539filesystem

Before getting started in implementing the proposed design in the previous section, let's define two new files: `str.h` and `str.c` which contain string related function that can be handy when we write our filesystem. Two functions will be implemented in `str.c` and they are `strcpy` which copies a string from a location to another, and `strcmp` which compares two strings, if they are equals then 1 is returned, otherwise 0 is returned. There is no need to explain the details of the code of these two functions since they depend on the normal way which C uses with strings. The following is the content of `str.c` and to save some space here I haven't provide `str.h` which only contains the prototypes of the two functions that have mentioned above.

```
#include "str.h"

void strcpy( char *dest, char *src )
{
    int idx = 0;

    while ( *src != '\0' )
    {
        dest[ idx ] = *src;

        src++;
        idx++;
    }
}

int strcmp( char *str1, char *str2 )
{
    while ( *str1 != '\0' )
    {
        if ( *str1 != *str2 )
            return 0;

        str1++;
        str2++;
    }
}
```

¹⁶In real-world situation, giving a whole block for a metadata of 260 bytes can be considered as a waste of space. One of real filesystems goals is to was as little space as possible to maintain the structure of the run-time filesystem.

```

    }

    if ( *str2 != '\0' )
        return 0;

    return 1;
}

```

Now we can start implementing 539filesystem, the first step as usual is to create the files that hold the functions related to our filesystem: `filesystem.h` and `filesystem.c`. The following is the content of `filesystem.h`.

```

#define BASE_BLOCK_ADDRESS 100
#define FILENAME_LENGTH 256

typedef struct
{
    int head, tail;
} base_block_t;

typedef struct
{
    char filename[ FILENAME_LENGTH ];
    int next_file_address;
} metadata_t;

base_block_t *base_block;

void filesystem_init();
int get_files_number();
void create_file( char *, char * );
char **list_files();
char *read_file( char * );

```

First we define two macros, `BASE_BLOCK_ADDRESS` and `FILENAME_LENGTH`. The first one is the address of base block in the disk, as we have mentioned earlier, this address is 100. The second one is the maximum length of a filename in 539filesystem, and we mentioned earlier that this length is 256. Then we define two structures as types: `base_block_t` and `metadata_t`, based on our discussions on 539filesystem design, you may have noticed that `base_block_t` represents the base block, it has two fields, each one of them of size 4 bytes, the first one is `head` and the second one is `tail`. The type `metadata_t` represents the metadata of a file, it has two fields as we described before, the first one is the filename and the second one is the metadata address of the next file. These two structures are based on linked-list data structure, and we are going to use them to load the data that they represent from the disk, manipulate them while they are in the main memory, then write them back to the disk in order to

make the run-time filesystem persistent. Then the global variable `base_block` is defined, which is the memory address that contains the base block after loading it from the disk, as we have said, this loaded copy is the one that we are going to implement when the user performs a transactional operation on the run-time filesystem such as creating a new file for example. The first function that we are going to implement is `filesystem_init` which is an initializer that will be called once the kernel starts. Its code is too simple, it is going to use the ATA device driver to read the base block from disk to the main memory and stores the memory address of this loaded data in the global variable `base_block`.

```
void filesystem_init()
{
    base_block = read_disk( BASE_BLOCK_ADDRESS );
}
```

We need to call this function by putting the line `filesystem_init();` in `kernel_main` of `main.c` after the line `scheduler_init();`. The rest of functions will be discussed in the following sub sections.

Creating a New File

Let's begin with the function `create_file` which is the one that create a new file. We mentioned before that there is no write operation in 539filesystem, instead, the content of a new file is written in the same operation that creates a new file. Basically, `create_file` operation should decide the disk address that the new file and its metadata should be stored in, of course, in real-world situation, the filesystem should be sure that this disk address is free and doesn't contain a part of another file. After deciding the disk address of this new file, we know that in the first block, the metadata of the file should be stored, and in the next block the content of this file should be stored, therefore, the metadata of the new file should be initialized by using the type `metadata_t` and after that should be stored into the disk by using ATA device driver right before the content of the file which is received by the function `create_file`. Beside writing the metadata and the content of the file in the disk, creating a new file in 539filesystem is equivalent to inserting a new item in a linked-list, that is, the base block need to be modified to make sure that the new file is reachable later. To do that, the metadata address of the new file should now be the tail in the base block, that is, the metadata address that was the tail before creating the new file is not the tail anymore, it become a normal item in the list that was once the tail, now, this previous tail should point to the newly created file in its next filed, and the tail in base block should be updated in the base block to point to the newly created file. There are also more subtle cases in updating the base block that will be discussed while writing the code of `create_file`. Let's start with the first part of the function.

```
void create_file( char *filename, char *buffer )
{
```

```

int metadata_lba = ( base_block->head == 0 ) ? BASE_BLOCK_ADDRESS + 1 : base_block->tail;
int file_lba = metadata_lba + 1;

metadata_t *metadata = kalloc( sizeof( metadata_t ) );

metadata->next_file_address = 0;

int currIdx;

for ( currIdx = 0; *filename != '\0' && currIdx < FILENAME_LENGTH - 1; currIdx++, filename++ )
    metadata->filename[ currIdx ] = *filename;

metadata->filename[ currIdx ] = '\0';

write_disk( metadata_lba, metadata );
write_disk( file_lba, buffer );

```

When the value of the head in the base block is 0, that means there is no files at all in the run-time filesystem. When `create_file` is called in this situation, that means this file that the caller is requesting to create is the first file in the run-time filesystem, the metadata of this first file can be simply stored in the block right after the base block. In `create_file` this fact is used to decide the disk address for the metadata of the new file, this address is stored in the local variable `metadata_lba` which its name is a short for “metadata logical block address”. In case that the run-time filesystem is not empty, that is, the value of `head` is not 0, then the tail field of base block can be used to decide the metadata disk address of the new file. As we know, the tail field contains the metadata address of the last file that has been added to the run-time filesystem, and the content of that file is stored in the disk address `tail + 1`, which means `tail + 2` is a free block that can be used to store new data ¹⁷, so we choose this address for the new metadata. After that, the disk address of the new content is decided by simply adding 1 to the disk address of the new metadata, the address of the content is stored in the local variable `file_lba`.

After deciding the disk addresses of the new metadata and file content, we start in creating the metadata of the file to store them later on the disk. As you can see, we allocate a space in the kernel’s heap for the new metadata by depending on the type `metadata_t`, after this allocation, we can use the local variable `metadata` to fill the fields of the new file metadata. First, we set the value of the next field to 0, because, as we mentioned earlier, the new file will be the tail file which means there is no file after it. Then, we copy the filename which is passed as a parameter `filename` to the filename field of the metadata, in case the passed filename’s length is less than the maximum length, then the whole filename is copied, otherwise, only the maximum number of characters of the

¹⁷This is ensured since 539filesystem stores the files in order, so, there will be no files after the tail unless it is a deleted file which can be overwritten and causes no data lose.

passed filename is copied and the rest are simply omitted. The final step that is related to the new file is to write the metadata and the file content on the right addresses in the disk, and this is done in the last two lines which use the ATA device driver to write the data to the disk. The following is the next and last part of `create_file` which updates the base block depending on the current state of the run-time filesystem.

```

    if ( base_block->head == 0 )
    {
        update_base_block( metadata_lba, metadata_lba );
    }
    else
    {
        metadata_t *tail_metadata = load_metadata( base_block->tail );

        tail_metadata->next_file_address = metadata_lba;

        write_disk( base_block->tail, tail_metadata );
        update_base_block( base_block->head, metadata_lba );
    }
} // End of "create_file"

```

When the run-time filesystem is empty, that is, the value of `head` in the base block is 0, then the new file that we are creating will be both the head and the tail file. As you can see, in the block of `if` statement that checks whether `head` equals 0, the not defined yet function `update_base_block` is used, the job of this function is updating the values of `head` and `tail` of the base block and write these changes on the permanent copy of the base block on the disk, the disk address of the new file's metadata is simply set as head and tail when `update_base_block` is called in this case. The second case is when the run-time filesystem isn't empty, that is, the value of `head` isn't 0. In this case we need to update the disk address of the tail in the base block to consider the new file as the new tail, furthermore, the next field of the previous tail, which is not the tail anymore, should be updated to point to the metadata of the new file, you see in `else` block that this is exactly what is done. The function that isn't defined yet `load_metadata` is used to load the metadata of the previous tail by passing the disk address of the metadata that we desire to load as a parameter. The local variable `tail_metadata` points to that loaded metadata of the tail, and depending on the type `metadata_t` we can reach the values of the previous tail fields easily. You can see that we simply change the value of the next field to the metadata address of the new file, then we write this modification on the disk and of course on the same location, finally, the tail file is updated in the base block by calling `update_base_block` which its code is presented next.

```

void update_base_block( int new_head, int new_tail )
{
    base_block->head = new_head;

```

```

    base_block->tail = new_tail;

    write_disk( BASE_BLOCK_ADDRESS, base_block );
}

```

It's too simple, it receives the value of head and tail that we would like to set on the base block, then, the copy of base block which is stored in the main memory is updated, then, this updated version is overwritten on the base block address on the disk. The following is code of `load_metadata` which has been used in `create_file` function.

```

metadata_t *load_metadata( int address )
{
    metadata_t *metadata = kalloc( sizeof( metadata_t ) );

    metadata = read_disk( address );

    return metadata;
}

```

Simply, it receives a disk address and assumes that the block which is presented by this address is a metadata block. It simply loads this metadata to the main memory by allocating a space for this metadata in the kernel's heap then loading the content of the address from the disk. The following is a sample of using `create_file` to create a new file in the run-time filesystem.

```

char *data = kalloc( 512 );

strcpy( data, "The content of the first file on 539filesystem" );

create_file( "first_file", data );

```